

Homework Help

```
# load libraries
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.3      v readr      2.1.4
v forcats    1.0.0      v stringr    1.5.0
v ggplot2    3.4.4      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.0
v purrr      1.0.2
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(mvtnorm)
library(coda)
library(rstanarm)
```

Loading required package: Rcpp

This is rstanarm version 2.26.1


- See <https://mc-stan.org/rstanarm/articles/priors> for changes to default priors!
- Default priors may change, so it's safest to specify priors, even if equivalent to the default
- For execution on a local, multicore CPU with excess RAM we recommend calling
`options(mc.cores = parallel::detectCores())`

```
yX = readRDS(
  url("http://www2.stat.duke.edu/~pdh10/Teaching/360/Materials/yXSS.rds"))
```

```
y<-yX[,1]
```

```
X<-yX[,-1]
```

```
# image(matrix(y,151,43))  
image(matrix(y,151,43))
```



$$y \sim N(X\beta, \sigma^2 I)$$

- y : 6493×1
- X : 6493×9
- β : 9×1
- σ^2 : 1×1
- I : 6493×6493

$$y_i = \beta_1 x_1 + \beta_2 x_2 + \dots \beta_p x_p + \epsilon_i$$

```
set.seed(360)  
# prior hyperparameters  
p = 9 # number of covariates  
Sigma0 = 1 * diag(rep(1, p)) # p dimension since no intercept.  
b0 = rep(1/9, p)  
nu0 = 2  
sigma02 = 1  
n = length(y)  
  
# starting values  
## note: gamma = 1 / sigma^2  
beta = t(rep(1/9, 9))
```

```

# values we should compute just once
SigmaInv = solve(Sigma0)
X2 = t(X) %*% X
Xy = t(X) %*% y
SIB0 = SigmaInv %*% b0
a = (nu0 + n) / 2
nu0s02 = nu0 * sigma02

## empty objects to fill
BETA = NULL
GAMMA = NULL

S = 2000
for (s in 1:S) {

  ### UPDATE SIGMA
  SSR1 = (y - (X %*% t(beta)))
  SSRB = t(SSR1) %*% SSR1
  gamma = rgamma(1, a, ((nu0s02 + SSRB) / 2))

  ### UPDATE BETA
  V = solve(SigmaInv + (gamma * X2))
  m = V %*% (Xy * gamma) # simplified since b0 = 0
  beta = rmvnorm(1, mean = m, sigma = V)

  ### SAVE STATES
  GAMMA = c(GAMMA, gamma)
  BETA = rbind(BETA, beta)
}

```

```
effectiveSize(BETA)
```

var1	var2	var3	var4	var5	var6	var7	var8
2000.000	2000.000	2000.000	2000.000	2000.000	2040.626	2000.000	2000.000
var9							
2000.000							

```
effectiveSize(GAMMA)
```

```
var1  
1810.087
```

```
# hist(1 / GAMMA)  
s2 = 1 / GAMMA  
sum(s2 > 100)
```

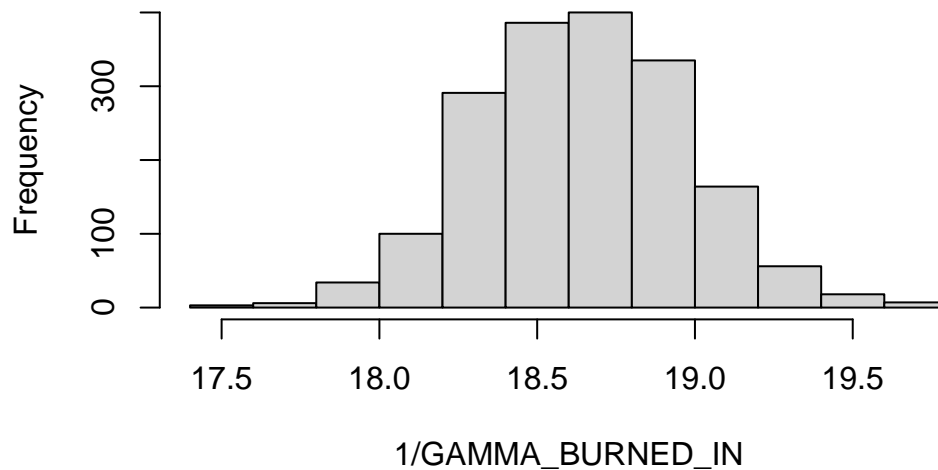
```
[1] 1
```

```
s2[1:5]
```

```
[1] 3086.86079 19.68145 18.75163 18.19946 18.33743
```

```
GAMMA_BURNED_IN = GAMMA[-c(1:S/10)]  
hist(1 / GAMMA_BURNED_IN)
```

Histogram of 1/GAMMA_BURNED_IN



```
BETA_BURNED = BETA[-c(1:S/10),]  
  
posteriorMeanBeta = apply(BETA_BURNED, 2, mean)  
posteriorCIBeta = apply(BETA_BURNED, 2, quantile, probs = c(0.025, 0.975))  
rbind(posteriorMeanBeta, posteriorCIBeta)
```

	[,1]	[,2]	[,3]	[,4]	[,5]
posteriorMeanBeta	0.5466149	-0.01868557	0.0002884716	0.049511233	-0.0122083365
2.5%	0.4300168	-0.05692517	-0.0163582932	0.004122091	-0.0255515678
97.5%	0.6608110	0.02140099	0.0168541549	0.096441935	0.0009236346

	[,6]	[,7]	[,8]	[,9]
posteriorMeanBeta	0.2741874	0.1606993	0.004890124	-0.08407046
2.5%	0.1969602	0.1133077	0.002942036	-0.22878357
97.5%	0.3453285	0.2089253	0.006761198	0.06937184

Predictors 1, 6, 7.

```
colnames(X)[c(1, 6, 7)]
```

```
[1] "Effluent" "Soil"      "Street"
```

Alternate method

- `rstanarm` uses software `stan`
- `stan` lets us sample from posteriors without writing a sampler ourselves each time.
- `stan` uses “HMC”: Hamiltonian Monte Carlo to sample from the posterior.
- one of the drawbacks of `stan` is that the specific prior specification we want may be difficult or impossible to obtain in the package.

```
yX_df = data.frame(yX)
post1 = stan_glm(V1 ~ 0 + ., data = yX_df, # remove the intercept
                 family = gaussian(link = "identity"),
                 seed = 360,
                 prior = normal(1/9, 1), # sets the beta prior
                 prior_aux = NULL) # set a flat prior on sigma
```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 6.6e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.66 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [20%] (Warmup)

```

Chain 1: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 1: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 1:
Chain 1: Elapsed Time: 0.75 seconds (Warm-up)
Chain 1: 0.993 seconds (Sampling)
Chain 1: 1.743 seconds (Total)
Chain 1:

```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).

```

Chain 2:
Chain 2: Gradient evaluation took 1.3e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.13 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration: 1 / 2000 [ 0%] (Warmup)
Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 1.109 seconds (Warm-up)
Chain 2: 1.072 seconds (Sampling)
Chain 2: 2.181 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 1.4e-05 seconds

```

Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration: 1 / 2000 [0%] (Warmup)
Chain 3: Iteration: 200 / 2000 [10%] (Warmup)
Chain 3: Iteration: 400 / 2000 [20%] (Warmup)
Chain 3: Iteration: 600 / 2000 [30%] (Warmup)
Chain 3: Iteration: 800 / 2000 [40%] (Warmup)
Chain 3: Iteration: 1000 / 2000 [50%] (Warmup)
Chain 3: Iteration: 1001 / 2000 [50%] (Sampling)
Chain 3: Iteration: 1200 / 2000 [60%] (Sampling)
Chain 3: Iteration: 1400 / 2000 [70%] (Sampling)
Chain 3: Iteration: 1600 / 2000 [80%] (Sampling)
Chain 3: Iteration: 1800 / 2000 [90%] (Sampling)
Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.862 seconds (Warm-up)
Chain 3: 1.101 seconds (Sampling)
Chain 3: 1.963 seconds (Total)
Chain 3:

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).

Chain 4:
Chain 4: Gradient evaluation took 3.8e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.38 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration: 1 / 2000 [0%] (Warmup)
Chain 4: Iteration: 200 / 2000 [10%] (Warmup)
Chain 4: Iteration: 400 / 2000 [20%] (Warmup)
Chain 4: Iteration: 600 / 2000 [30%] (Warmup)
Chain 4: Iteration: 800 / 2000 [40%] (Warmup)
Chain 4: Iteration: 1000 / 2000 [50%] (Warmup)
Chain 4: Iteration: 1001 / 2000 [50%] (Sampling)
Chain 4: Iteration: 1200 / 2000 [60%] (Sampling)
Chain 4: Iteration: 1400 / 2000 [70%] (Sampling)
Chain 4: Iteration: 1600 / 2000 [80%] (Sampling)
Chain 4: Iteration: 1800 / 2000 [90%] (Sampling)
Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 1.309 seconds (Warm-up)

```
Chain 4:          1.013 seconds (Sampling)
Chain 4:          2.322 seconds (Total)
Chain 4:
```

```
cbind(post1$coefficients, posteriorMeanBeta)
```

		posteriorMeanBeta
Effluent	0.5489818902	0.5466148619
Influent	-0.0193440302	-0.0186855712
Poultry	0.0004494069	0.0002884716
Reference	0.0511734579	0.0495112329
Septic	-0.0117853543	-0.0122083365
Soil	0.2742320104	0.2741873716
Street	0.1588910894	0.1606993061
Swine	0.0048560389	0.0048901240
Wetland	-0.0855007899	-0.0840704585

```
mean(sqrt(1 / GAMMA_BURNED_IN))
```

```
[1] 4.317466
```